

Ordered hash map: search tree optimized by a hash table

Petar Ivanov
Valentina Dyankova
Biserka Yovcheva

Shumen University, Bulgaria

Use case: sorted databases with **few updates vs many finds**

Widely used approaches:

- **Hash map** – fast finds but **no ordering!**
- **Balanced search tree (BST)** – ordered but **slow finds!**

Solution:

wrap **BST** + **hash map**

Binary search tree (BST) for ordered data

Hash map for fast finds

Optimizations

- Keys **only** in BST (hash map of pointers)
- Values **only** in BST (hash map of pointers)
- Fast erase – marking as erased

Complexity

Asymptotic estimate	map	ordered hash map	ordered hash map* (fast erase)	unordered map
insert	logN	logN	logN*	1
find	logN	1	1	1
erase	logN	logN	1	1
next/prev traversal	logN	logN	logN*	n/a or N
ordered traversal	N	N	N*	n/a or NlogN

N* – # of ever inserted elements

Experiments

[sec]	map	ordered hash map	ordered hash map*	unorderedmap
insert	2.68	3.60	3.63	0.83
find	2.67	0.51	0.63	0.49
erase	3.32	4.05	0.94	0.91
next/prev traversal	0.21	0.22	0.94	n/a

Applications

- Sensitive hashing and other locality-sensitive hashes
- All STL map/set uses for hashable elements

Conclusion

Open source C++ classes and experiments at
<https://github.com/petar-ivanov/ordered-hash>